

▼ Telco Customer Churn Prediction Report

Author: Thattaphol Puttawithee

Tools: Python (Pandas, Scikit-learn, XGBoost), Power BI, SQL Server

Dataset: Telco Customer Churn (Kaggle)

Date: October 2025

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

▼ Import data to Jupyter Notebook

Query data from SQL Server

```
import pandas as pd
from sqlalchemy import create_engine

pd.set_option('display.max_columns', None)

engine = create_engine(
    r"mssql+pyodbc://LAPTOP-QC1AHOCH\SQLEXPRESS/telco_db"
    r"?driver=ODBC+Driver+17+for+SQL+Server&trusted_connection=yes"
)

df = pd.read_sql("SELECT * FROM stg_Churn;", engine)
df
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	DeviceProtection	TechSupport	StreamingTV	StreamingMovies	Contract	PaperlessBilling	PaymentMethod	MonthlyCharges	TotalCharges
0	0002-ORFBO	Female	0	Yes	Yes	9	Yes	No	DSL	No	Yes	No	Yes	Yes	No	One year	Yes	Mailed check	65.599998	593.2
1	0003-MKNFE	Male	0	No	No	9	Yes	Yes	DSL	No	No	No	No	No	Yes	Month-to-month	No	Mailed check	59.900002	542.4
2	0004-TLHLJ	Male	0	No	No	4	Yes	No	Fiber optic	No	No	Yes	No	No	No	Month-to-month	Yes	Electronic check	73.900002	280.8
3	0011-IGKFF	Male	1	Yes	No	13	Yes	No	Fiber optic	No	Yes	Yes	No	Yes	Yes	Month-to-month	Yes	Electronic check	98.000000	1237.8
4	0013-EXCHZ	Female	1	Yes	No	3	Yes	No	Fiber optic	No	No	No	Yes	Yes	No	Month-to-month	Yes	Mailed check	83.900002	267.3
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
7038	9987-LUTYD	Female	0	No	No	13	Yes	No	DSL	Yes	No	No	Yes	No	No	One year	No	Mailed check	55.150002	742.9
7039	9992-RRAMN	Male	0	Yes	No	22	Yes	Yes	Fiber optic	No	No	No	No	No	Yes	Month-to-month	Yes	Electronic check	85.099998	1873.6
7040	9992-UJOEL	Male	0	No	No	2	Yes	No	DSL	No	Yes	No	No	No	No	Month-to-month	Yes	Mailed check	50.299999	92.7
7041	9993-LHIEB	Male	0	Yes	Yes	67	Yes	No	DSL	Yes	No	Yes	Yes	No	Yes	Two year	No	Mailed check	67.849998	4627.6
7042	9995-HOTOH	Male	0	Yes	Yes	63	No	No phone service	DSL	Yes	Yes	Yes	No	Yes	Yes	Two year	No	Electronic check	59.000000	3707.6

7043 rows × 21 columns

```
df["ChargeRatio"] = df["TotalCharges"] / (df["tenure"] + 1)
df
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	DeviceProtection	TechSupport	StreamingTV	StreamingMovies	Contract	PaperlessBilling	PaymentMethod	MonthlyCharges	TotalCharges
0	0002-ORFBO	Female	0	Yes	Yes	9	Yes	No	DSL	No	Yes	No	Yes	Yes	No	One year	Yes	Mailed check	65.599998	593.200544
1	0003-MKNFE	Male	0	No	No	9	Yes	Yes	DSL	No	No	No	No	No	Yes	Month-to-month	No	Mailed check	59.900002	542.400352
2	0004-TLHLJ	Male	0	No	No	4	Yes	No	Fiber optic	No	No	Yes	No	No	No	Month-to-month	Yes	Electronic check	73.900002	280.800576
3	0011-IGKFF	Male	1	Yes	No	13	Yes	No	Fiber optic	No	Yes	Yes	No	Yes	Yes	Month-to-month	Yes	Electronic check	98.000000	1237.800672
4	0013-EXCHZ	Female	1	Yes	No	3	Yes	No	Fiber optic	No	No	No	Yes	Yes	No	Month-to-month	Yes	Mailed check	83.900002	267.300384
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
7038	9987-LUTYD	Female	0	No	No	13	Yes	No	DSL	Yes	No	No	Yes	No	No	One year	No	Mailed check	55.150002	742.900256
7039	9992-RRAMN	Male	0	Yes	No	22	Yes	Yes	Fiber optic	No	No	No	No	No	Yes	Month-to-month	Yes	Electronic check	85.099998	1873.600448
7040	9992-UJOEL	Male	0	No	No	2	Yes	No	DSL	No	Yes	No	No	No	No	Month-to-month	Yes	Mailed check	50.299999	92.700096
7041	9993-LHIEB	Male	0	Yes	Yes	67	Yes	No	DSL	Yes	No	Yes	Yes	No	Yes	Two year	No	Mailed check	67.849998	4627.600256
7042	9995-HOTOH	Male	0	Yes	Yes	63	No	No phone service	DSL	Yes	Yes	Yes	No	Yes	Yes	Two year	No	Electronic check	59.000000	3707.600384

7043 rows × 22 columns

```
df["PaymentMethod"].unique()
```

```
array(['Mailed check', 'Electronic check', 'Credit card (automatic)',
      'Bank transfer (automatic)'], dtype=object)
```

```
df["PaymentAutomatic"] = df["PaymentMethod"].apply(lambda x: 1 if "automatic" in x.lower() else 0)
```

```
for m in [12, 24, 36, 48, 60]:
    df[f"tenure_ge_{m}m"] = (df["tenure"] >= m).astype(int)
```

```
df.info()
df.describe()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 28 columns):
#   Column              Non-Null Count  Dtype
---  -
0   customerID          7043 non-null   object
1   gender              7043 non-null   object
2   SeniorCitizen        7043 non-null   object
3   Partner              7043 non-null   object
4   Dependents          7043 non-null   object
5   tenure              7043 non-null   int64
6   PhoneService        7043 non-null   object
7   MultipleLines        7043 non-null   object
8   InternetService      7043 non-null   object
9   OnlineSecurity       7043 non-null   object
10  OnlineBackup         7043 non-null   object
11  DeviceProtection     7043 non-null   object
12  TechSupport          7043 non-null   object
13  StreamingTV          7043 non-null   object
14  StreamingMovies      7043 non-null   object
15  Contract             7043 non-null   object
16  PaperlessBilling     7043 non-null   object
17  PaymentMethod        7043 non-null   object
18  MonthlyCharges       7043 non-null   float64
19  TotalCharges         7032 non-null   float64
20  Churn                7043 non-null   object
21  ChargeRatio          7032 non-null   float64
22  PaymentAutomatic     7043 non-null   int64
23  tenure_ge_12m        7043 non-null   int64
24  tenure_ge_24m        7043 non-null   int64
25  tenure_ge_36m        7043 non-null   int64
26  tenure_ge_48m        7043 non-null   int64
27  tenure_ge_60m        7043 non-null   int64
dtypes: float64(3), int64(7), object(18)
memory usage: 1.5+ MB
```

	tenure	MonthlyCharges	TotalCharges	ChargeRatio	PaymentAutomatic	tenure_ge_12m	tenure_ge_24m	tenure_ge_36m	tenure_ge_48m	tenure_ge_60m
count	7043.000000	7043.000000	7032.000000	7032.000000	7043.000000	7043.000000	7043.000000	7043.000000	7043.000000	7043.000000
mean	32.371149	64.761692	2283.300441	59.083067	0.435326	0.706233	0.557575	0.433196	0.326991	0.210564
std	24.559481	30.090047	2266.771363	30.514438	0.495835	0.455519	0.496709	0.495552	0.469147	0.407738
min	0.000000	18.250000	18.799999	9.183333	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	9.000000	35.500000	401.449997	26.225944	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
50%	29.000000	70.349998	1397.475037	61.070387	0.000000	1.000000	1.000000	0.000000	0.000000	0.000000
75%	55.000000	89.849998	3794.737488	84.877538	1.000000	1.000000	1.000000	1.000000	1.000000	0.000000
max	72.000000	118.750000	8684.799805	118.969860	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

▼ Check NULL

Check null value in the data

```
df.isna().sum()
```

```
customerID    0
gender         0
SeniorCitizen 0
Partner        0
Dependents     0
tenure         0
PhoneService   0
MultipleLines  0
InternetService 0
OnlineSecurity 0
OnlineBackup   0
DeviceProtection 0
TechSupport    0
StreamingTV    0
StreamingMovies 0
Contract       0
PaperlessBilling 0
PaymentMethod  0
MonthlyCharges 0
TotalCharges   11
Churn          0
ChargeRatio    11
PaymentAutomatic 0
tenure_ge_12m 0
tenure_ge_24m 0
tenure_ge_36m 0
tenure_ge_48m 0
tenure_ge_60m 0
dtype: int64
```


Remove null

```
df = df.dropna()
```

Check Outliers


```
sns.boxplot(x=df['tenure'])
```

<Axes: xlabel='tenure'>



```
sns.scatterplot(x='tenure', y='MonthlyCharges', hue='Churn', data=df)
```


<Axes: xlabel='tenure', ylabel='MonthlyCharges'>



Check data balance

```
df['Churn'].value_counts().plot(kind='bar')
```

<Axes: xlabel='Churn'>



df

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	DeviceProtection	TechSupport	StreamingTV	StreamingMovies	Contract	PaperlessBilling	PaymentMethod	MonthlyCharges	TotalCh
0	0002-ORFBO	Female	0	Yes	Yes	9	Yes	No	DSL	No	Yes	No	Yes	Yes	No	One year	Yes	Mailed check	65.599998	593.2
1	0003-MKNFE	Male	0	No	No	9	Yes	Yes	DSL	No	No	No	No	No	Yes	Month-to-month	No	Mailed check	59.900002	542.4
2	0004-TLHLJ	Male	0	No	No	4	Yes	No	Fiber optic	No	No	Yes	No	No	No	Month-to-month	Yes	Electronic check	73.900002	280.8
3	0011-IGKFF	Male	1	Yes	No	13	Yes	No	Fiber optic	No	Yes	Yes	No	Yes	Yes	Month-to-month	Yes	Electronic check	98.000000	1237.8
4	0013-EXCHZ	Female	1	Yes	No	3	Yes	No	Fiber optic	No	No	No	Yes	Yes	No	Month-to-month	Yes	Mailed check	83.900002	267.3
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
7038	9987-LUTYD	Female	0	No	No	13	Yes	No	DSL	Yes	No	No	Yes	No	No	One year	No	Mailed check	55.150002	742.9
7039	9992-RRAMN	Male	0	Yes	No	22	Yes	Yes	Fiber optic	No	No	No	No	No	Yes	Month-to-month	Yes	Electronic check	85.099998	1873.6
7040	9992-UJOEL	Male	0	No	No	2	Yes	No	DSL	No	Yes	No	No	No	No	Month-to-month	Yes	Mailed check	50.299999	92.7
7041	9993-LHIEB	Male	0	Yes	Yes	67	Yes	No	DSL	Yes	No	Yes	Yes	No	Yes	Two year	No	Mailed check	67.849998	4627.6
7042	9995-HOTOH	Male	0	Yes	Yes	63	No	No phone service	DSL	Yes	Yes	Yes	No	Yes	Yes	Two year	No	Electronic check	59.000000	3707.6

7032 rows × 28 columns

```
binary_cols = ['Partner', 'Dependents', 'PhoneService', 'MultipleLines',
               'OnlineSecurity', 'OnlineBackup', 'DeviceProtection',
               'TechSupport', 'StreamingTV', 'StreamingMovies',
               'PaperlessBilling', 'Churn']
```

```
df[binary_cols] = df[binary_cols].replace({'Yes': 1, 'No': 0})
```

```
df['gender'] = df['gender'].replace({'Male': 1, 'Female': 0})
df['MultipleLines'] = df['MultipleLines'].replace({'No phone service': -1})
df
```

C:\Temp\ipykernel\_3172\2260154420.py:6: FutureWarning: Downcasting behavior in `replace` is deprecated and will be removed in a future version. To retain the old behavior, explicitly call `result.infer\_objects(copy=False)`. To opt-in to the future behavior, set `pd.set\_option('future.no\_silent\_downcasting', True)`  
df[binary\_cols] = df[binary\_cols].replace({'Yes': 1, 'No': 0})  
C:\Temp\ipykernel\_3172\2260154420.py:6: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df[binary_cols] = df[binary_cols].replace({'Yes': 1, 'No': 0})
C:\Temp\ipykernel_3172\2260154420.py:7: FutureWarning: Downcasting behavior in `replace` is deprecated and will be removed in a future version. To retain the old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-in to the future behavior, set `pd.set_option('future.no_silent_downcasting', True)`
df['gender'] = df['gender'].replace({'Male': 1, 'Female': 0})
C:\Temp\ipykernel_3172\2260154420.py:7: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df['gender'] = df['gender'].replace({'Male': 1, 'Female': 0})
C:\Temp\ipykernel_3172\2260154420.py:8: FutureWarning: Downcasting behavior in `replace` is deprecated and will be removed in a future version. To retain the old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-in to the future behavior, set `pd.set_option('future.no_silent_downcasting', True)`
df['MultipleLines'] = df['MultipleLines'].replace({'No phone service': -1})
C:\Temp\ipykernel_3172\2260154420.py:8: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df['MultipleLines'] = df['MultipleLines'].replace({'No phone service': -1})
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	DeviceProtection	TechSupport	StreamingTV	StreamingMovies	Contract	PaperlessBilling	PaymentMethod	MonthlyCharges	TotalCharges
0	0002-ORFBO	0	0	1	1	9	1	0	DSL	0	1	0	1	1	0	One year	1	Mailed check	65.599998	593.2
1	0003-MKNFE	1	0	0	0	9	1	1	DSL	0	0	0	0	0	1	Month-to-month	0	Mailed check	59.900002	542.4
2	0004-TLHLJ	1	0	0	0	4	1	0	Fiber optic	0	0	1	0	0	0	Month-to-month	1	Electronic check	73.900002	280.8
3	0011-IGKFF	1	1	1	0	13	1	0	Fiber optic	0	1	1	0	1	1	Month-to-month	1	Electronic check	98.000000	1237.8
4	0013-EXCHZ	0	1	1	0	3	1	0	Fiber optic	0	0	0	1	1	0	Month-to-month	1	Mailed check	83.900002	267.3
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
7038	9987-LUTYD	0	0	0	0	13	1	0	DSL	1	0	0	1	0	0	One year	0	Mailed check	55.150002	742.9
7039	9992-RRAMN	1	0	1	0	22	1	1	Fiber optic	0	0	0	0	0	1	Month-to-month	1	Electronic check	85.099998	1873.6
7040	9992-UJOEL	1	0	0	0	2	1	0	DSL	0	1	0	0	0	0	Month-to-month	1	Mailed check	50.299999	92.7
7041	9993-LHIEB	1	0	1	1	67	1	0	DSL	1	0	1	1	0	1	Two year	0	Mailed check	67.849998	4627.6
7042	9995-HOTOH	1	0	1	1	63	0	-1	DSL	1	1	1	0	1	1	Two year	0	Electronic check	59.000000	3707.6
7032 rows × 28 columns																				

7032 rows × 28 columns

```
Services_cols = (('OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV', 'StreamingMovies'))
df[Services_cols] = df[Services_cols].replace({'No internet service': 0})
df.isin(['No internet service']).sum()
```

C:\Temp\ipykernel\_3172\3187367540.py:2: FutureWarning: Downcasting behavior in `replace` is deprecated and will be removed in a future version. To retain the old behavior, explicitly call `result.infer\_objects(copy=False)`. To opt-in to the future behavior, set `pd.set\_option('future.no\_silent\_downcasting', True)`  
df[Services\_cols] = df[Services\_cols].replace({'No internet service': 0})  
C:\Temp\ipykernel\_3172\3187367540.py:2: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df[Services_cols] = df[Services_cols].replace({'No internet service': 0})
```

```
customerID      0
gender          0
SeniorCitizen   0
Partner         0
Dependents      0
tenure          0
PhoneService    0
MultipleLines   0
InternetService  0
OnlineSecurity  0
OnlineBackup    0
DeviceProtection 0
TechSupport     0
StreamingTV     0
StreamingMovies 0
Contract        0
PaperlessBilling 0
PaymentMethod   0
MonthlyCharges  0
TotalCharges    0
Churn           0
ChargeRatio     0
PaymentAutomatic 0
tenure_ge_12m   0
tenure_ge_24m   0
tenure_ge_36m   0
tenure_ge_48m   0
tenure_ge_60m   0
dtype: int64
```

```
df['SeniorCitizen'] = df['SeniorCitizen'].astype(int)
```

C:\Temp\ipykernel\_3172\2584085859.py:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df['SeniorCitizen'] = df['SeniorCitizen'].astype(int)
```

```
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
numeric_cols = ['tenure', 'MonthlyCharges', 'TotalCharges', 'ChargeRatio']
```

```
df[numeric_cols] = scaler.fit_transform(df[numeric_cols])
df
```

C:\Temp\ipykernel\_3172\3296849014.py:6: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df[numeric_cols] = scaler.fit_transform(df[numeric_cols])
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	DeviceProtection	TechSupport	StreamingTV	StreamingMovies	Contract	PaperlessBilling	PaymentMethod	MonthlyCharges	Total	
0	0002-ORFBO	0	0	1	1	-0.954296		1	0	DSL	0	1	0	1	1	0	One year	1	Mailed check	0.026652	-
1	0003-MKNFE	1		0	0	-0.954296		1	1	DSL	0	0	0	0	0	1	Month-to-month	0	Mailed check	-0.162819	-
2	0004-TLHLJ	1		0	0	-1.158016		1	0	Fiber optic	0	0	1	0	0	0	Month-to-month	1	Electronic check	0.302548	-
3	0011-IGKFF	1		1	1	-0.791321		1	0	Fiber optic	0	1	1	0	1	1	Month-to-month	1	Electronic check	1.103642	-
4	0013-EXCHZ	0		1	1	-1.198760		1	0	Fiber optic	0	0	0	1	1	0	Month-to-month	1	Mailed check	0.634952	-
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	
7038	9987-LUTYD	0		0	0	-0.791321		1	0	DSL	1	0	0	1	0	0	One year	0	Mailed check	-0.320711	-
7039	9992-RRAMN	1		0	1	-0.424625		1	1	Fiber optic	0	0	0	0	0	1	Month-to-month	1	Electronic check	0.674841	-
7040	9992-UJOEL	1		0	0	-1.239504		1	0	DSL	0	1	0	0	0	0	Month-to-month	1	Mailed check	-0.481927	-
7041	9993-LHIEB	1		0	1	1.408853		1	0	DSL	1	0	1	1	0	1	Two year	0	Mailed check	0.101443	-
7042	9995-HOTOH	1		0	1	1.245878		0	-1	DSL	1	1	1	0	1	1	Two year	0	Electronic check	-0.192735	-

7032 rows × 28 columns

```
categorical_cols = ['InternetService', 'Contract', 'PaymentMethod']
```

```
df = pd.get_dummies(df, columns=categorical_cols, drop_first=True, dtype=int)
df
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	OnlineSecurity	OnlineBackup	DeviceProtection	TechSupport	StreamingTV	StreamingMovies	PaperlessBilling	MonthlyCharges	TotalCharges	Churn	ChargeRatio	PaymentAutor
	0	0002-ORFBO	0	0	1	1	-0.954296	1	0	0	1	0	1	1	0	1	0.026652	-0.745607	0	0.008093
	1	0003-MKNFE	1	0	0	0	-0.954296	1	1	0	0	0	0	0	1	0	-0.162819	-0.768063	0	-0.158725
	2	0004-TLHLJ	1	0	0	0	-1.158016	1	0	0	0	1	0	0	0	1	0.302548	-0.883456	1	-0.095472
	3	0011-IGKFF	1	1	1	0	-0.791321	1	0	0	1	1	0	1	1	1	1.103642	-0.461240	1	0.961410
	4	0013-EXCHZ	0	1	1	0	-1.198760	1	0	0	0	0	1	1	0	1	0.634952	-0.889390	1	0.254551
	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
	7038	9987-LUTYD	0	0	0	0	-0.791321	1	0	1	0	0	1	0	0	0	-0.320711	-0.679605	0	-0.197258
	7039	9992-RRAMN	1	0	1	0	-0.424625	1	1	0	0	0	0	0	1	1	0.674841	-0.180711	1	0.733546
	7040	9992-UJOEL	1	0	0	0	-1.239504	1	0	0	1	0	0	0	0	1	-0.481927	-0.966443	0	-0.923117
	7041	9993-LHIEB	1	0	1	1	1.408853	1	0	1	0	1	1	0	1	0	0.101443	1.034298	0	0.294000
	7042	9995-HOTOH	1	0	1	1	1.245878	0	-1	1	1	1	0	1	1	0	-0.192735	0.628383	0	-0.037749

7032 rows × 32 columns

```
df.info()

<class 'pandas.core.frame.DataFrame'>
Index: 7032 entries, 0 to 7042
Data columns (total 32 columns):
#   Column                                Non-Null Count  Dtype
---  ---                                -
0   customerID                            7032 non-null   object
1   gender                                7032 non-null   int64
2   SeniorCitizen                         7032 non-null   int64
3   Partner                               7032 non-null   int64
4   Dependents                           7032 non-null   int64
5   tenure                               7032 non-null   float64
6   PhoneService                         7032 non-null   int64
7   MultipleLines                        7032 non-null   int64
8   OnlineSecurity                       7032 non-null   int64
9   OnlineBackup                         7032 non-null   int64
10  DeviceProtection                     7032 non-null   int64
11  TechSupport                          7032 non-null   int64
12  StreamingTV                          7032 non-null   int64
13  StreamingMovies                      7032 non-null   int64
14  PaperlessBilling                     7032 non-null   int64
15  MonthlyCharges                       7032 non-null   float64
16  TotalCharges                         7032 non-null   float64
17  Churn                                7032 non-null   int64
18  ChargeRatio                          7032 non-null   float64
19  PaymentAutomatic                     7032 non-null   int64
20  tenure_ge_12m                       7032 non-null   int64
21  tenure_ge_24m                       7032 non-null   int64
22  tenure_ge_36m                       7032 non-null   int64
23  tenure_ge_48m                       7032 non-null   int64
24  tenure_ge_60m                       7032 non-null   int64
25  InternetService_Fiber optic         7032 non-null   int64
26  InternetService_No                   7032 non-null   int64
27  Contract_One year                    7032 non-null   int64
28  Contract_Two year                    7032 non-null   int64
29  PaymentMethod_Credit card (automatic) 7032 non-null   int64
30  PaymentMethod_Electronic check      7032 non-null   int64
31  PaymentMethod_Mailed check          7032 non-null   int64
dtypes: float64(4), int64(27), object(1)
memory usage: 1.8+ MB
```

Train Test split


```
from sklearn.model_selection import train_test_split
X = df.drop(['Churn', 'customerID', 'TotalCharges', 'gender'], axis=1)
y = df['Churn']

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

X_train.shape, X_test.shape, y_train.shape, y_test.shape

((5625, 28), (1407, 28), (5625,), (1407,))
```

Class imbalance



Prepare SMOTE

```
import pandas as pd
from imblearn.over_sampling import SMOTE

sm = SMOTE(random_state=42)
X_train_res, y_train_res = sm.fit_resample(X_train, y_train)

y_train_series = pd.Series(y_train)
y_train_res_series = pd.Series(y_train_res)

print("Before SMOTE:", y_train_series.value_counts().to_dict())
print("After SMOTE :", y_train_res_series.value_counts().to_dict())

Before SMOTE: {0: 4130, 1: 1495}
After SMOTE : {0: 4130, 1: 4130}
```

Logistic Regression

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, roc_auc_score

model_lr = LogisticRegression(class_weight='balanced', max_iter=1000) # balanced for imbalanced data
model_lr.fit(X_train, y_train)
y_pred_lr = model_lr.predict(X_test)

print(classification_report(y_test, y_pred_lr))
print("ROC AUC:", roc_auc_score(y_test, model_lr.predict_proba(X_test)[:,:1]))

precision  recall  f1-score  support
0         0.90    0.74    0.82    1033
1         0.53    0.78    0.63     374

accuracy           0.75    1407
macro avg         0.71    0.76    0.72    1407
weighted avg      0.80    0.75    0.77    1407

ROC AUC: 0.8509261224510926
```

RandomForest

```
from sklearn.ensemble import RandomForestClassifier

model_rf = RandomForestClassifier(
    n_estimators=300,
    class_weight='balanced_subsample', # balanced_subsample for imbalanced data
    max_depth=10,
    random_state=42
)
```

```
model_rf.fit(X_train, y_train)
rf_pred = model_rf.predict(X_test)

print("ROC AUC:", roc_auc_score(y_test, model_rf.predict_proba(X_test)[:,-1]))
print(classification_report(y_test, rf_pred))
```

```
ROC AUC: 0.8445405883905969
      precision    recall  f1-score   support

      0       0.89       0.79       0.84       1033
      1       0.56       0.72       0.63        374

 accuracy          0.78       1407
macro avg          0.72       0.76       0.74       1407
weighted avg          0.80       0.78       0.78       1407
```

▼ XGB

```
from xgboost import XGBClassifier

xgb = XGBClassifier(
    scale_pos_weight=(len(y_train[y_train==0]) / len(y_train[y_train==1])),
    learning_rate=0.05,
    n_estimators=500,
    max_depth=4,
    subsample=0.8,
    colsample_bytree=0.8,
    random_state=42
)
xgb.fit(X_train, y_train)
xgb_pred = xgb.predict(X_test)

print("ROC AUC:", roc_auc_score(y_test, xgb.predict_proba(X_test)[:,-1]))
print("\nClassification Report:\n", classification_report(y_test, xgb_pred))
```

```
ROC AUC: 0.8408806187264134

Classification Report:
      precision    recall  f1-score   support

      0       0.90       0.76       0.82       1033
      1       0.53       0.75       0.62        374

 accuracy          0.76       1407
macro avg          0.71       0.76       0.72       1407
weighted avg          0.80       0.76       0.77       1407
```

▼ Hyperparameter Tuning

Logistic Regression

```
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import RandomizedSearchCV, StratifiedKFold
from scipy.stats import loguniform

cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

lr = LogisticRegression(class_weight='balanced', solver='liblinear', max_iter=5000)

lr_param = {
    "C": loguniform(1e-3, 1e2), # 0.001..100
    "penalty": ["l1", "l2"],
}

lr_rs = RandomizedSearchCV(
    lr, lr_param, n_iter=40, scoring="roc_auc", cv=cv,
    n_jobs=-1, random_state=42, error_score="raise"
)
lr_rs.fit(X_train, y_train)
print("Best parameters:", lr_rs.best_params_)
```

```
Best parameters: {'C': np.float64(0.5414413211338525), 'penalty': 'l2'}
```

```
import numpy as np
from sklearn.metrics import precision_score, recall_score, f1_score, roc_auc_score, classification_report

best_lr = lr_rs.best_estimator_
y_proba = best_lr.predict_proba(X_test)[:,-1]

for t in [0.3, 0.4, 0.5, 0.6, 0.7]:
    y_pred_t = (y_proba >= t).astype(int)
    precision = precision_score(y_test, y_pred_t)
    recall = recall_score(y_test, y_pred_t)
    f1 = f1_score(y_test, y_pred_t)
    print(f"Threshold = {t:.1f} | Precision = {precision:.3f} | Recall = {recall:.3f} | F1 = {f1:.3f}")
```

```
Threshold = 0.3 | Precision = 0.429 | Recall = 0.922 | F1 = 0.586
Threshold = 0.4 | Precision = 0.478 | Recall = 0.869 | F1 = 0.617
Threshold = 0.5 | Precision = 0.524 | Recall = 0.783 | F1 = 0.628
Threshold = 0.6 | Precision = 0.575 | Recall = 0.690 | F1 = 0.627
Threshold = 0.7 | Precision = 0.650 | Recall = 0.545 | F1 = 0.593
```

XMG

```
from xgboost import XGBClassifier
from sklearn.model_selection import RandomizedSearchCV, StratifiedKFold
from scipy.stats import randint, uniform, loguniform
from sklearn.metrics import roc_auc_score

cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

# For XGBoost class imbalance
scale = len(y_train[y_train==0]) / len(y_train[y_train==1] * 0.8)

xgb = XGBClassifier(
    objective='binary:logistic',
    eval_metric='auc',
    scale_pos_weight=scale,
    random_state=42
)

xgb_param = {
    "n_estimators": randint(300, 800),
    "learning_rate": uniform(0.01, 0.2),
    "max_depth": randint(3, 7),
    "min_child_weight": randint(1, 8),
    "subsample": uniform(0.6, 0.4),
    "colsample_bytree": uniform(0.6, 0.4),
    "gamma": uniform(0, 0.5),
    "reg_lambda": loguniform(1e-2, 1e2),
    "reg_alpha": loguniform(1e-3, 1e1),
}

xgb_rs = RandomizedSearchCV(
    estimator=xgb,
    param_distributions=xgb_param,
    n_iter=50,
    scoring='roc_auc',
    cv=cv,
    n_jobs=-1,
    random_state=42,
    verbose=1
)

xgb_rs.fit(X_train, y_train)

print("Best parameters:", xgb_rs.best_params_)
```

```
Fitting 5 folds for each of 50 candidates, totalling 250 fits
Best parameters: {'colsample_bytree': np.float64(0.8232408008069365), 'gamma': np.float64(0.2019180855290204), 'learning_rate': np.float64(0.02297844942179631), 'max_depth': 4, 'min_child_weight': 4, 'n_estimators': 567, 'reg_alpha': np.float64(0.010165510266418737), 'reg_lambda': np.float64(0.9749762207436125), 'subsample':
```

```
import numpy as np
from sklearn.metrics import precision_score, recall_score, f1_score, roc_auc_score, classification_report

best_xgb = xgb_rs.best_estimator_
y_proba = best_xgb.predict_proba(X_test)[:,-1]

for t in [0.3, 0.4, 0.5, 0.6, 0.7]:
    y_pred_t = (y_proba >= t).astype(int)
    precision = precision_score(y_test, y_pred_t)
    recall = recall_score(y_test, y_pred_t)
    f1 = f1_score(y_test, y_pred_t)
    auc = roc_auc_score(y_test, y_proba)
    print(f"Threshold = {t:.1f} | Precision = {precision:.3f} | Recall = {recall:.3f} | F1 = {f1:.3f} | {auc:<10.3f}")
```


Threshold = 0.3 | Precision = 0.445 | Recall = 0.909 | F1 = 0.598 | 0.848  
Threshold = 0.4 | Precision = 0.499 | Recall = 0.866 | F1 = 0.633 | 0.848  
Threshold = 0.5 | Precision = 0.527 | Recall = 0.786 | F1 = 0.631 | 0.848  
Threshold = 0.6 | Precision = 0.573 | Recall = 0.693 | F1 = 0.627 | 0.848  
Threshold = 0.7 | Precision = 0.640 | Recall = 0.551 | F1 = 0.592 | 0.848

CrossValidation (threshold = 0.4)

```
import matplotlib.pyplot as plt
from sklearn.metrics import precision_recall_curve

precision, recall, thresholds = precision_recall_curve(y_test, y_proba)

plt.figure(figsize=(6,5))
plt.plot(thresholds, precision[:-1], label='Precision', color='blue')
plt.plot(thresholds, recall[:-1], label='Recall', color='orange')
plt.xlabel('Threshold')
plt.ylabel('Score')
plt.title('Precision vs Recall Tradeoff')
plt.legend()
plt.grid(True)
plt.show()
```



```
from sklearn.model_selection import StratifiedKFold
from sklearn.metrics import precision_score, recall_score, f1_score, roc_auc_score
import numpy as np

cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

precisions, recalls, f1s, aucs = [], [], [], []

for train_idx, val_idx in cv.split(X_train, y_train):
    X_tr, X_val = X_train.iloc[train_idx], X_train.iloc[val_idx]
    y_tr, y_val = y_train.iloc[train_idx], y_train.iloc[val_idx]

    # Fit model
    model = best_xgb
    model.fit(X_tr, y_tr)

    # Predict probability
    y_proba = model.predict_proba(X_val)[:, 1]

    # Apply threshold = 0.4
    y_pred = (y_proba >= 0.4).astype(int)

    # Calculate metrics
    precisions.append(precision_score(y_val, y_pred))
    recalls.append(recall_score(y_val, y_pred))
    f1s.append(f1_score(y_val, y_pred))
    aucs.append(roc_auc_score(y_val, y_proba))

print("Threshold-based Cross Validation (t=0.4)")
print(f"ROC-AUC : {np.mean(aucs):.3f} ± {np.std(aucs):.3f}")
print(f"Precision : {np.mean(precisions):.3f} ± {np.std(precisions):.3f}")
print(f"Recall : {np.mean(recalls):.3f} ± {np.std(recalls):.3f}")
print(f"F1 : {np.mean(f1s):.3f} ± {np.std(f1s):.3f}")
```

Threshold-based Cross Validation (t=0.4)  
ROC-AUC : 0.842 ± 0.009  
Precision : 0.487 ± 0.007  
Recall : 0.840 ± 0.022  
F1 : 0.616 ± 0.008

```
# FINAL MODEL - XGBoost (threshold = 0.4)
import joblib
import numpy as np

final_xgb = xgb_rs.best_estimator_ # ㄱ RandomizedSearchCV
final_threshold = 0.4


# save model and threshold
joblib.dump(final_xgb, "final_xgb_model.joblib")
np.save("final_threshold.npy", np.array([final_threshold]))
print("Model and threshold saved.")
```

Model and threshold saved.

```
import pandas as pd
import matplotlib.pyplot as plt

importance = pd.DataFrame({
    "Feature": X_train.columns,
    "Importance": final_xgb.feature_importances_
}).sort_values("Importance", ascending=False)

plt.figure(figsize=(8,4))
plt.barh(importance["Feature"], importance["Importance"])
plt.gca().invert_yaxis()
plt.title("Top 10 Important Features (XGBoost)")
plt.show()
```



## Executive Summary

“Developed a customer churn prediction model using XGBoost, achieving 0.84 ROC-AUC. At threshold = 0.4, the model achieves 0.50 Precision, 0.86 Recall, and 0.63 F1-score. Insights show that contract type, tenure, and monthly charges are the strongest predictors of churn.”

### Key Insights

- Long-term contracts → drastically reduce churn
- Fiber optic customers → high churn risk
- Low tenure + high monthly charges → most likely to churn

### Recommended Business Actions

- Offer loyalty discounts or incentives to short-tenure customers to reduce early churn.
- Improve satisfaction for fiber-optic customers through proactive service monitoring and feedback collection.
- Retarget paperless billing users with higher satisfaction campaigns